Overview Cost & Tokens Tool Usage User Behavior Operational Health Advanced Analytics

Overview

High-level metrics across all Claude Code usage

Total Cost	Sessions	Active Users	API Calls	Error Rate
$6,001.43	5,000	100	118,014	1.15%

Daily User Activity

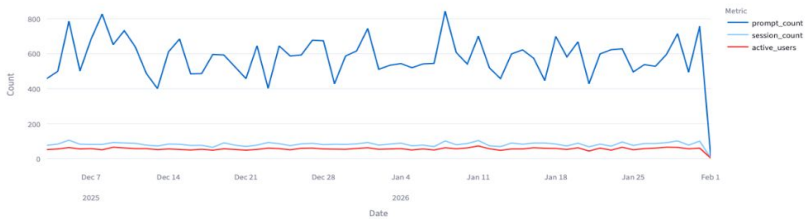

Daily Cost Trend

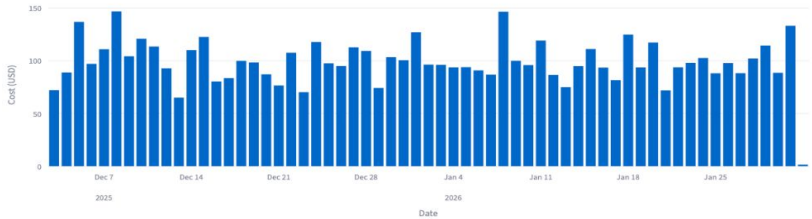

Data Quality

Fields Monitored	Parse Failures	Fields with Issues
12	0	0

	table_name	field_name	total_rows	null_count	null_pct	parse_failure_count	parse_fail_pct	is_optional
0	api_errors	attempt	1362	0	0	0	0	☐
1	api_errors	duration_ms	1362	0	0	0	0	☐
2	api_requests	cache_creation_tokens	118014	0	0	0	0	☐
3	api_requests	cache_read_tokens	118014	0	0	0	0	☐
4	api_requests	cost_usd	118014	0	0	0	0	☐
5	api_requests	duration_ms	118014	0	0	0	0	☐
6	api_requests	input_tokens	118014	0	0	0	0	☐
7	api_requests	output_tokens	118014	0	0	0	0	☐
8	tool_results	duration_ms	148418	0	0	0	0	☐
9	tool_results	result_size_bytes	148418	104067	70.12	0	0	☑

1. Overview

$6,001 in total cost across 5,000 sessions, 100 unique users, and 61 days, with a 1.15% API error rate and zero detected data quality issues.
The org averages about 82 sessions per day. All 100 employees were active at least once during the analysis window, but average daily participation was about 56 users. The average session costs $1.20 and contains 7.03 turns. The busiest day was 2025-12-05, with 106 sessions from 64 users.

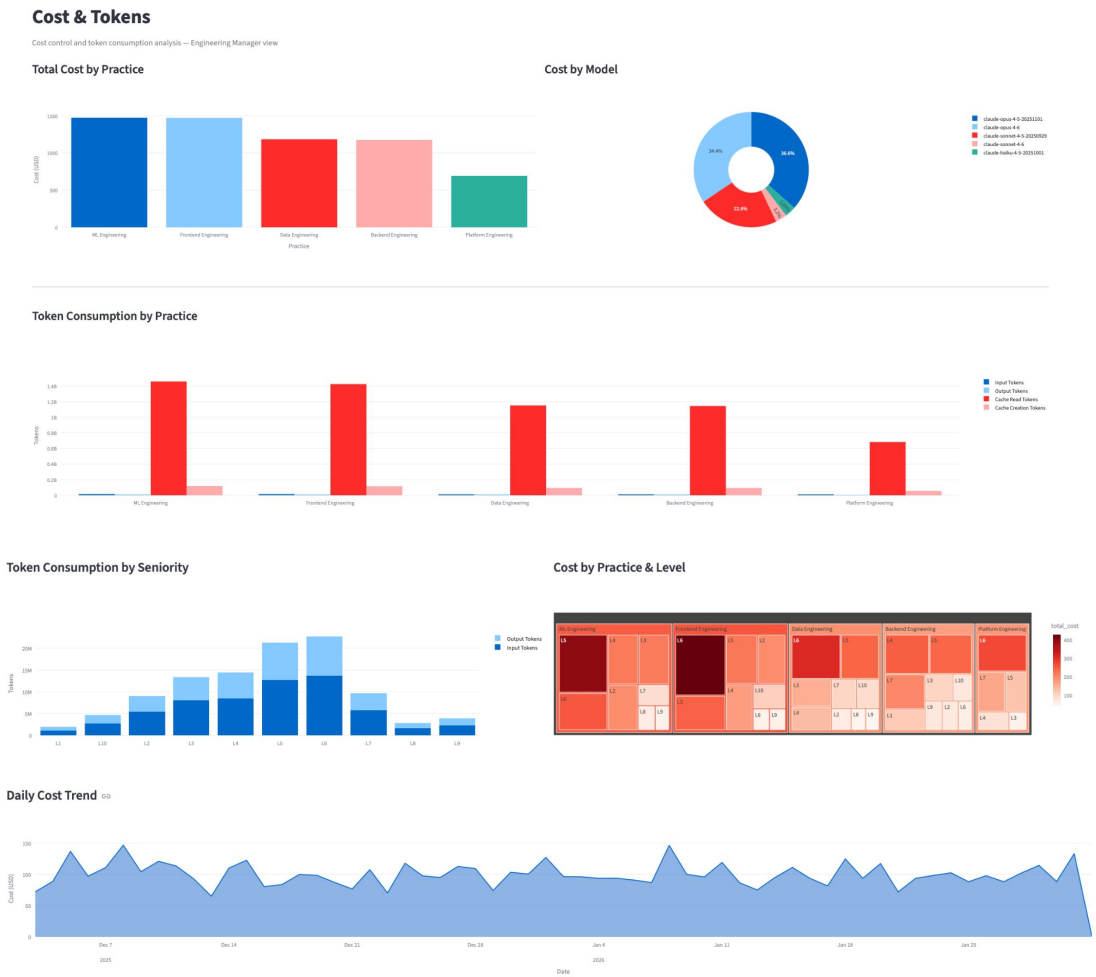

2. Cost & Tokens

91.0% of processed tokens were cache reads, showing that prompt caching is the dominant cost-control mechanism.
ML Engineering ($1,475) and Frontend Engineering ($1,473) are the top spenders, but cost per user is relatively flat across practices at about $53-$62. Opus requests cost roughly $0.08-$0.09 on average versus $0.004 for Haiku, yet Haiku still handles 39.0% of all requests (45,979), which suggests clear model tiering for lightweight tasks. Cost per user by seniority is also fairly uniform, ranging from $50.93 to $69.62, with L10 highest.

3. Tool Usage

Read (30.4%) and Bash (28.5%) account for 58.9% of all tool decisions, with Edit a distant third at 12.6%.

Bash has the lowest success rate at 93.0%, making it the clearest source of tool-level friction. A few smaller tools also sit below 97% success, notably WebSearch (93.6%), WebFetch (95.0%), and `mcp_tool` (97.0%, just under the cutoff). Rejection rates are flat across seniority at roughly 1.7%-2.2%, so there is no meaningful trust gap by level.

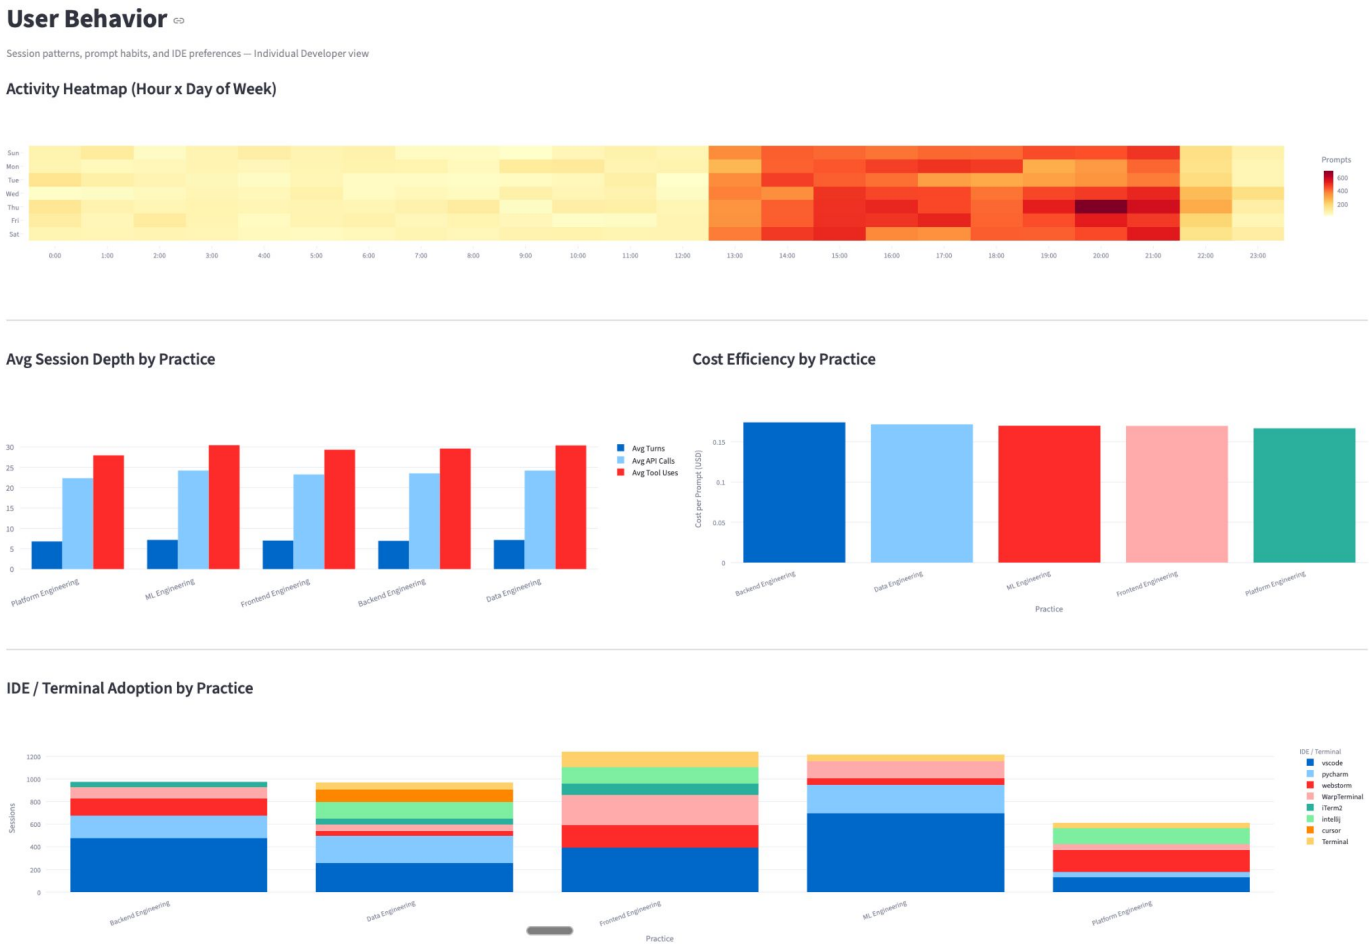

4. User Behavior

Usage peaks in the afternoon and evening, with the heaviest activity from 15:00-21:00 and 21:00 as the single busiest hour.
VSCode dominates IDE usage with 39 distinct users (39%), followed by PyCharm (15) and WebStorm (13). The top individual spender is blake.kim (Data Engineering, L10) at $94.44 across 58 sessions, or 1.57% of total org cost. L10 engineers have the deepest and most expensive sessions on average at 8.10 turns and $1.43 per session, while L4 has the shallowest sessions at 6.45 turns. Error rates are highest at the edges of the ladder, with L1 narrowly above L10 on average errors per session.

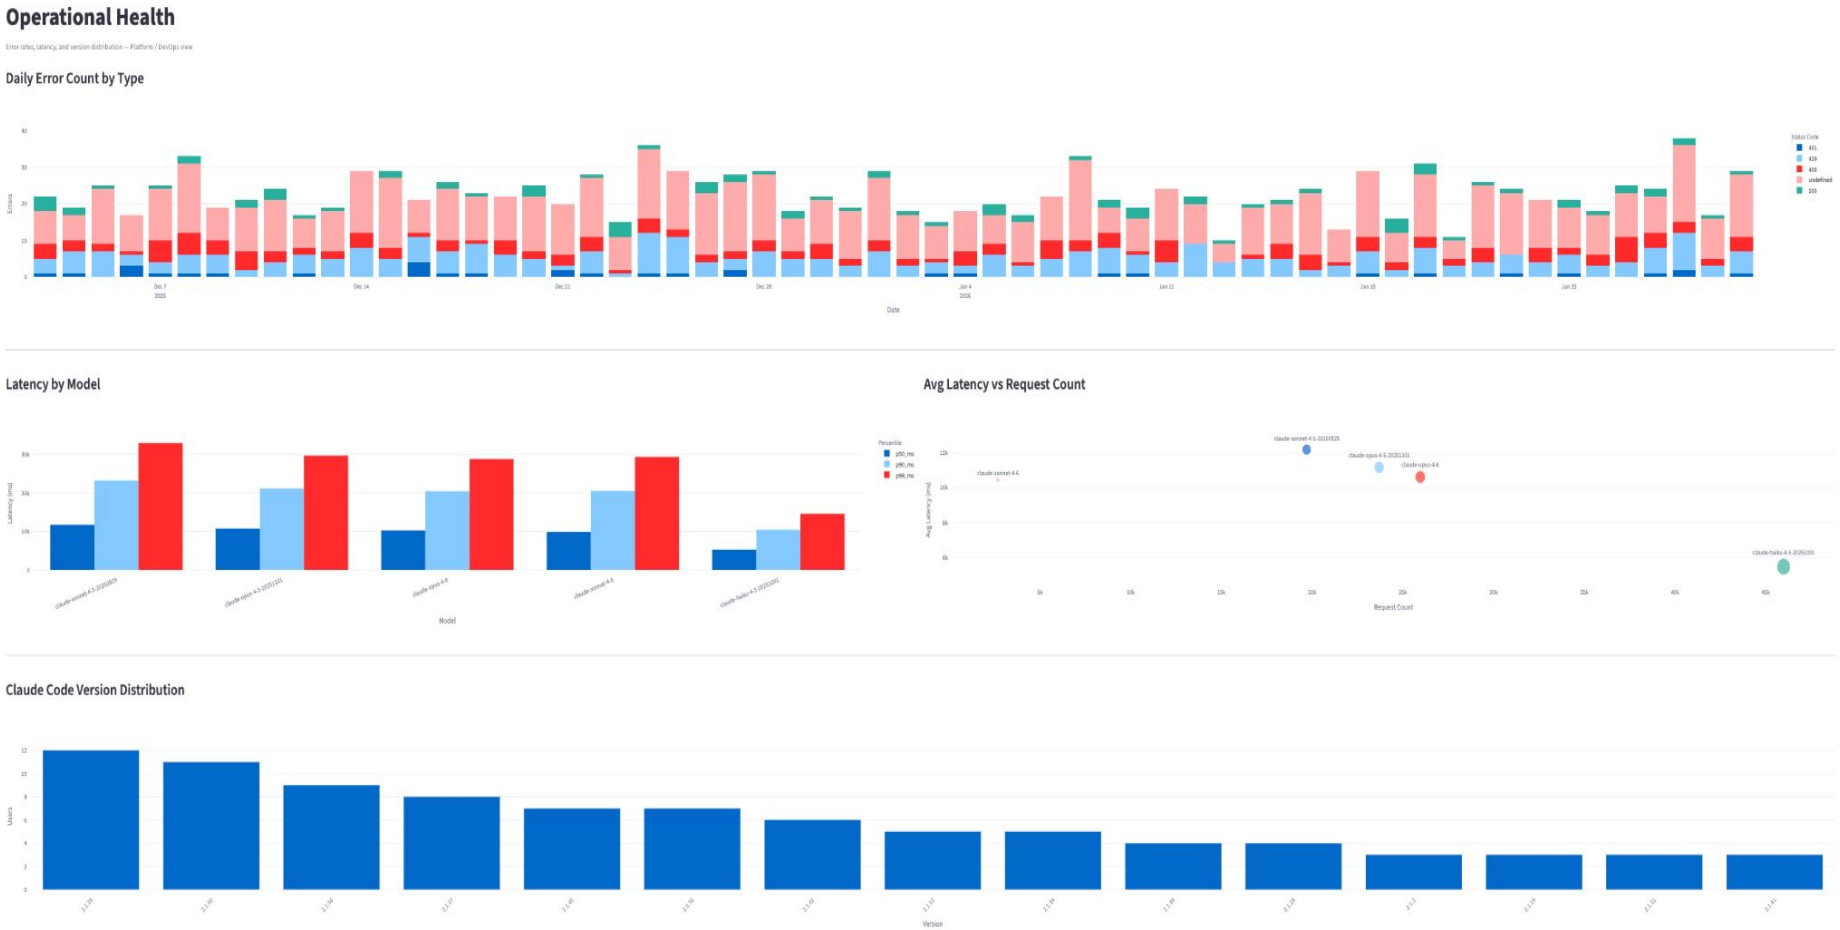

5. Operational Health

"Request was aborted." is the top error category with 717 of 1,362 total errors, affecting all 100 users. Rate limits (429s) are second at 299 occurrences affecting 95 users. Together, those two categories account for 74.6% of all errors.

Model error rates are tightly clustered from 0.82% to 1.24%, so no model is materially less reliable than the others. Haiku is about 2x faster than Opus and Sonnet on average at 5.48s versus roughly 10.4-12.2s, while Sonnet 4.5 has the highest p95 latency at 26.62s. 56.8% of all errors have status `undefined`, which suggests many are client-side aborts rather than server-side failures.

Advanced Analytics

Predictive analytics, anomaly detection, and user clustering

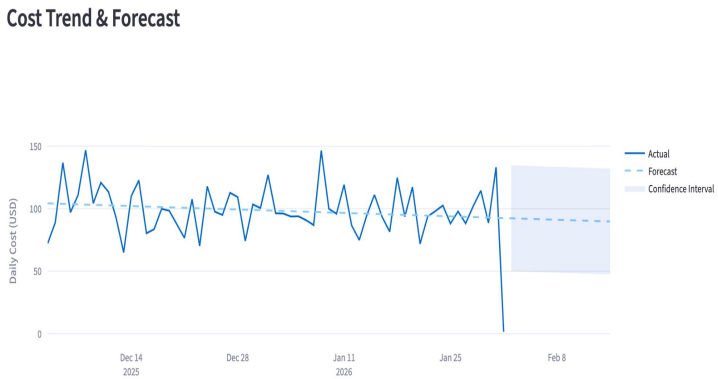

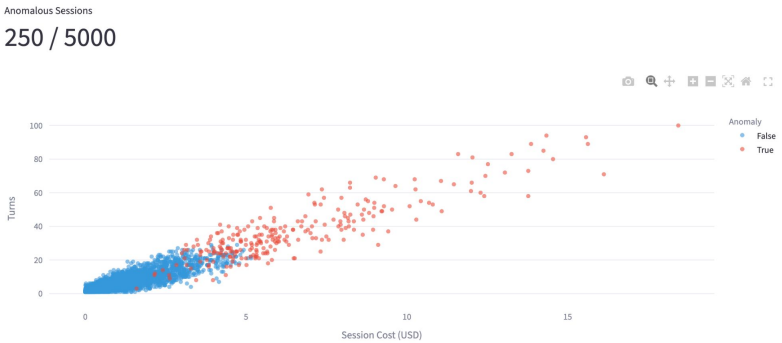

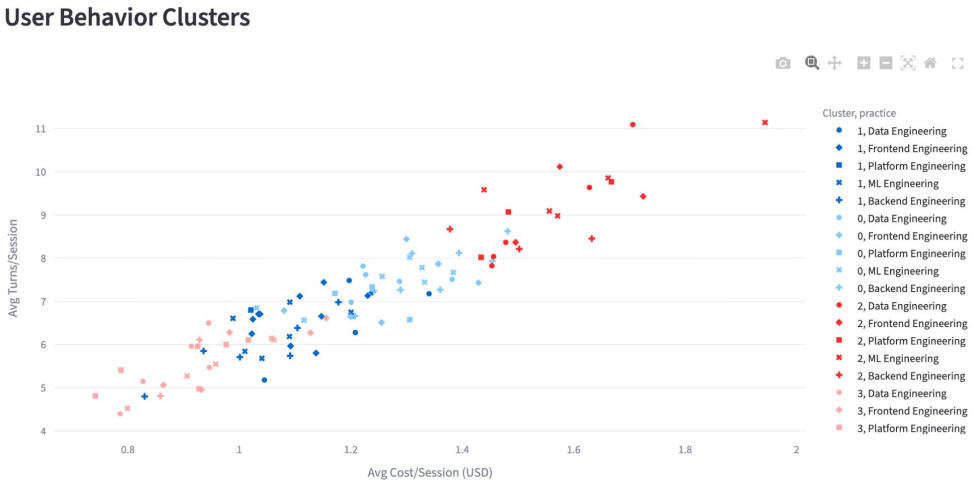

6. Advanced Analytics

Session cost is strongly right-skewed: median $0.67, mean $1.20, p95 $4.20, and max $18.44. Most sessions are cheap, while a small tail drives the average upward.

Weekly cost is stable at roughly $507-$752 with 99-100 active users per week, so there is no obvious adoption ramp or churn. Canada leads total spend at $1,549 from 26 users, while the United States is lowest at $524 from 10 users, largely because of headcount. At the practice aggregate level, behavior is remarkably uniform: all practices average about 6.8-7.2 turns, 27.9-30.4 tool uses, and $1.13-$1.23 per session.

7. Final Takeaways

Adoption is broad and stable, not a rollout. Weekly active users stayed at 99-100 across the full 61-day window, so Claude Code looks less like a pilot and more like established engineering infrastructure. Spend is also not concentrated in a tiny set of champions: the top individual user accounts for just 1.57% of total cost, and the top 10 users together account for 14.4%.

Cost appears naturally bounded. Per-user spend is remarkably flat across practices ($53-$62) and seniority levels ($51-$70), while Haiku handles 39% of all requests and Opus remains the premium model for heavier work. Just as important, 91% of processed tokens are cache reads, which suggests caching is doing most of the cost-control work.

Seniority changes session depth more than approval behavior. Rejection rates stay flat across levels at roughly 1.7%-2.2%, but session depth and cost rise at the senior end: L10 engineers average 8.10 turns and $1.43 per session. Error rates per session are also elevated at the extremes, with L10 among the highest and L1 slightly higher still.

Reliability noise is dominated by client-side and quota-related issues, not model-specific instability. `Request was aborted` accounts for 52.6% of all errors, rate limits add another 22.0%, and true 500-level server failures are only 6.1%. Model error rates are tightly clustered between 0.82% and 1.24%, so no single model stands out as materially less reliable.

Claude Code is primarily an afternoon and evening tool. Usage peaks between 15:00 and 21:00, with 21:00 as the single busiest hour, which aligns more with deep-work blocks and end-of-day pushes than morning planning time.

Cost outliers are rare but meaningful. The median session costs $0.67, but the p95 reaches $4.20 and the maximum hits $18.44, so a small tail of expensive sessions pulls the average above the median. Those outliers are worth monitoring, but whether they represent waste or valuable complex work is ultimately a management judgment, not a purely statistical one.